

Credit Card Approval

Import Libraries

```
In [85]: import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import numpy as np
from sklearn.neighbors import KNeighborsClassifier
from sklearn.model_selection import train_test_split, cross_val_score, KFold
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier, VotingClassifier
import xgboost as xgb
from sklearn.preprocessing import StandardScaler
```

Load Dataset

```
In [86]: #load the dataset
df = pd.read_csv('credit_risk_dataset.csv')

# Display the first few rows of the dataset
df.head()
```

```
Out [86]:   person_age  person_income  person_home_ownership  person_emp_length  loan_intent  loan_amnt  loan_int_rate  loan_status  loan_percent_income
0         22        59000          RENT                 123.0    PERSONAL    35000     16.02        1       0.59
1         21        9600           OWN                  5.0    EDUCATION    1000     11.14        0       0.10
2         25        9600        MORTGAGE                 1.0    MEDICAL     5500     12.87        1       0.57
3         23       65500          RENT                  4.0    MEDICAL    35000     15.23        1       0.53
4         24       54400          RENT                  8.0    MEDICAL    35000     14.27        1       0.55
```

Dataset Overview

```
In [87]: # Check the size of the dataset
df.shape
```

```
Out [87]: (32581, 11)
```

```
In [88]: # Display basic information about the dataset
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 32581 entries, 0 to 32580
Data columns (total 11 columns):
 #   Column            Non-Null Count  Dtype  
 --- 
 0   person_age        32581 non-null   int64  
 1   person_income      32581 non-null   int64  
 2   person_home_ownership  32581 non-null   object  
 3   person_emp_length  31686 non-null   float64 
 4   loan_intent        32581 non-null   object  
 5   loan_amnt          32581 non-null   int64  
 6   loan_int_rate      29465 non-null   float64 
 7   loan_status         32581 non-null   int64  
 8   loan_percent_income 32581 non-null   float64 
 9   cb_person_default_on_file 32581 non-null   object  
 10  cb_person_cred_hist_length 32581 non-null   int64  
dtypes: float64(3), int64(5), object(3)
memory usage: 2.7+ MB
```

Check and Handle Missing Values

```
In [89]: # Check for missing values in the dataset
df.isnull().sum()
```

```
Out [89]: person_age          0
person_income        0
person_home_ownership 0
person_emp_length     895
loan_intent          0
loan_amnt            0
```

```
loan_int_rate      3116
loan_status        0
loan_percent_income 0
cb_person_default_on_file 0
cb_person_cred_hist_length 0
dtype: int64
```

In [90]:

```
#check if missing values are indicated as nan
column = ['person_age', 'person_income', 'person_home_ownership', 'person_emp_length', 'loan_intent', 'loan_amnt', 'loan_int_rate', 'loan_status', 'loan_percent_income', 'cb_person_default_on_file', 'cb_person_cred_hist_length']

for missing_nan in column:
    if df[missing_nan].isna().any():
        print(f'Missing values are written as nan in column {missing_nan}')
```

```
Missing values are written as nan in column person_emp_length
Missing values are written as nan in column loan_int_rate
```

In [91]:

```
# Verify that missing values are handled
df.isnull().sum()
```

```
Out [91]: person_age          0
person_income         0
person_home_ownership 0
person_emp_length     895
loan_intent           0
loan_amnt             0
loan_int_rate         3116
loan_status           0
loan_percent_income   0
cb_person_default_on_file 0
cb_person_cred_hist_length 0
dtype: int64
```

In [92]:

```
# Replace NaN values in 'person_emp_length' with 0
df['person_emp_length'] = df['person_emp_length'].fillna(0)
```

In [93]:

```
from sklearn.impute import SimpleImputer
impute = SimpleImputer(strategy='median')
df['loan_int_rate'] = impute.fit_transform(df[['loan_int_rate']])
```

In [94]:

```
# Impute missing target values with the mean
#loan_int_rate_mean = df['loan_int_rate'].median()
#df['loan_int_rate'].fillna(loan_int_rate_mean, inplace=True)
```

Encoding Categorical Values

In [95]:

```
df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 32581 entries, 0 to 32580
Data columns (total 11 columns):
 #   Column            Non-Null Count  Dtype  
 --- 
 0   person_age        32581 non-null   int64  
 1   person_income      32581 non-null   int64  
 2   person_home_ownership 32581 non-null   object 
 3   person_emp_length  32581 non-null   float64
 4   loan_intent        32581 non-null   object 
 5   loan_amnt          32581 non-null   int64  
 6   loan_int_rate       32581 non-null   float64
 7   loan_status         32581 non-null   int64  
 8   loan_percent_income 32581 non-null   float64
 9   cb_person_default_on_file 32581 non-null   object 
 10  cb_person_cred_hist_length 32581 non-null   int64  
dtypes: float64(3), int64(5), object(3)
memory usage: 2.7+ MB
```

From the dataset info, person_home_ownership, loan_intent & cb_person_default_on_file columns are categorical values

In [96]:

```
#check for categorical unique values
df['person_home_ownership'].unique()
```

```
Out [96]: array(['RENT', 'OWN', 'MORTGAGE', 'OTHER'], dtype=object)
```

In [97]:

```
#create a dictionary to map the categorical value to numeric value equivalent
ownership = {'RENT':1, 'OWN':2, 'MORTGAGE':3, 'OTHER':4}

## Replace the categorical values in 'person_home_ownership' with their numeric equivalents
df['person_home_ownership'] = df['person_home_ownership'].replace(ownership)
```

```
In [98]: #check if the change has been effected  
df.head()
```

```
Out [98]:   person_age  person_income  person_home_ownership  person_emp_length  loan_intent  loan_amnt  loan_int_rate  loan_status  loan_percent_income  
0  22          59000            1                  123.0    PERSONAL     35000      16.02        1       0.59  
1  21          9600             2                   5.0    EDUCATION     1000      11.14        0       0.10  
2  25          9600             3                   1.0    MEDICAL      5500      12.87        1       0.57  
3  23          65500            1                  4.0    MEDICAL     35000      15.23        1       0.53  
4  24          54400            1                  8.0    MEDICAL     35000      14.27        1       0.55
```

```
In [99]: #check for loan_intent categorical column unique value  
df['loan_intent'].unique()
```

```
Out [99]: array(['PERSONAL', 'EDUCATION', 'MEDICAL', 'VENTURE', 'HOMEIMPROVEMENT',  
       'DEBTCONSOLIDATION'], dtype=object)
```

```
In [100]:  
#create a dictionary to map the categorical value to numeric value equivalent  
loan_intent = {'PERSONAL':1, 'EDUCATION':2, 'MEDICAL':3, 'VENTURE':4, 'HOMEIMPROVEMENT':5,  
              'DEBTCONSOLIDATION':6}  
  
## Replace the categorical values in 'loan_intent' with their numeric equivalents  
df['loan_intent'] = df['loan_intent'].replace(loan_intent)
```

```
In [101]: #check if the change has been effected, display dataset  
df.head()
```

```
Out [101]:   person_age  person_income  person_home_ownership  person_emp_length  loan_intent  loan_amnt  loan_int_rate  loan_status  loan_percent_income  
0  22          59000            1                  123.0        1     35000      16.02        1       0.59  
1  21          9600             2                   5.0        2     1000      11.14        0       0.10  
2  25          9600             3                   1.0        3      5500      12.87        1       0.57  
3  23          65500            1                  4.0        3     35000      15.23        1       0.53  
4  24          54400            1                  8.0        3     35000      14.27        1       0.55
```

```
In [102]: #create a dictionary to map the categorical value to numeric value equivalent  
cb = {'Y':1, 'N':0}  
  
#Replace the categorical values in 'loan_intent' with their numeric equivalents  
df['cb_person_default_on_file'] = df['cb_person_default_on_file'].replace(cb)
```

```
In [103]: #check if the change has been effected, display dataset  
df.head()
```

```
Out [103]:   person_age  person_income  person_home_ownership  person_emp_length  loan_intent  loan_amnt  loan_int_rate  loan_status  loan_percent_income  
0  22          59000            1                  123.0        1     35000      16.02        1       0.59  
1  21          9600             2                   5.0        2     1000      11.14        0       0.10  
2  25          9600             3                   1.0        3      5500      12.87        1       0.57  
3  23          65500            1                  4.0        3     35000      15.23        1       0.53  
4  24          54400            1                  8.0        3     35000      14.27        1       0.55
```

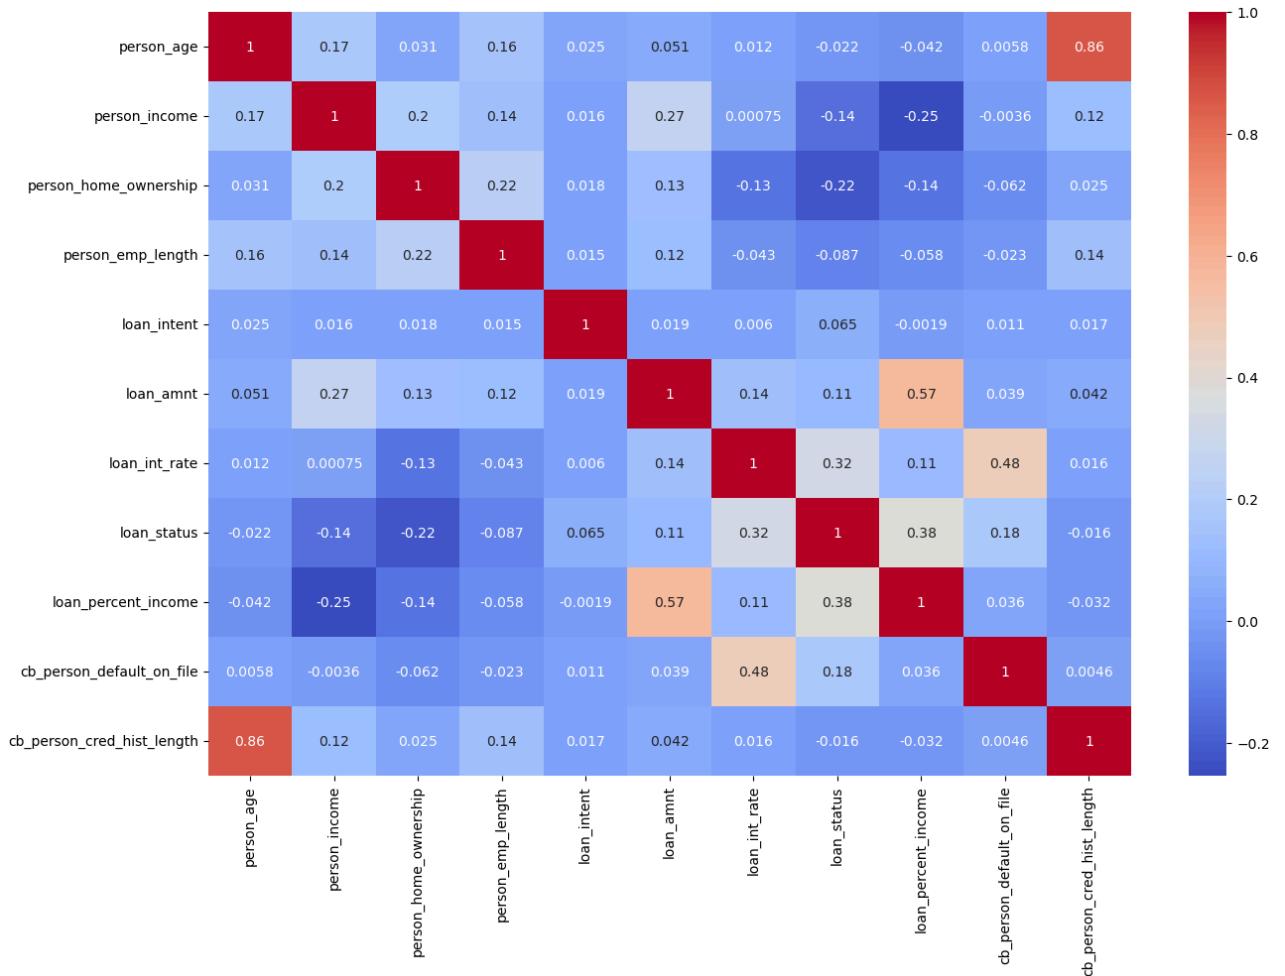
```
In [104]: df.isnull().sum()
```

```
Out [104]: person_age      0  
person_income      0  
person_home_ownership      0  
person_emp_length      0  
loan_intent      0  
loan_amnt      0  
loan_int_rate      0  
loan_status      0  
loan_percent_income      0  
cb_person_default_on_file      0  
cb_person_cred_hist_length      0  
dtype: int64
```

Correlation Matrix

```
In [105]: import seaborn as sns  
import matplotlib.pyplot as plt
```

```
# Check for linear relationships using a correlation matrix
plt.figure(figsize=(15,10))
correlation_matrix = df.corr()
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm')
plt.show()
```



Save and Reload Transformed Dataset

```
In [106]: # Save the transformed file as a CSV file
df.to_csv('credit_risk.csv', index=False)
```

```
In [107]: # Reload the transformed dataset
df = pd.read_csv('credit_risk.csv')
```

```
In [108]: #Display few rows of transformed dataset
df.head()
```

	person_age	person_income	person_home_ownership	person_emp_length	loan_intent	loan_amnt	loan_int_rate	loan_status	loan_percent_income	cb_person_default_on_file	cb_person_cred_hist_length
0	22	59000	1	123.0	1	35000	16.02	1	0.59		
1	21	9600	2	5.0	2	1000	11.14	0	0.10		
2	25	9600	3	1.0	3	5500	12.87	1	0.57		
3	23	65500	1	4.0	3	35000	15.23	1	0.53		
4	24	54400	1	8.0	3	35000	14.27	1	0.55		

Feature and Target Variables

```
In [109]: #Define feature and target variable
X = df.drop('loan_status', axis=1).values
y = df['loan_status'].values
```

Train-Test Split

```
In [110]: # Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X,y, test_size=0.2, random_state=21)
```

Data Standardization

```
In [111]: # Standardize the feature variables
scale = StandardScaler()
scale_train = scale.fit_transform(X_train)
scale_test = scale.transform(X_test)
```

Initialize Classifiers

```
In [112]: # Initialize classifiers
from xgboost import XGBClassifier
lg = LogisticRegression()
knn = KNeighborsClassifier()
#xgb = XGBClassifier()
#rnn = RandomForestClassifier()

#Create a list of classifiers for the Voting Classifier
classifiers = [('Logistic Regression', lg), ('KNeighbor Classifier', knn)]# ('XGBoost', xgb), ('Random Forest', rnn)
```

Voting Classifier

```
In [113]: #instantite Voting Classifier, the algorithm that defines the final predictions from all the models

vc = VotingClassifier(estimators = classifiers, voting='soft')

#fit model
vc.fit(scale_train, y_train)

#make final predictions
vpred = vc.predict(scale_test)

# Calculate accuracy for the individual classifiers
accuracies = {}

for name, model in classifiers:
    # Fit each classifier on the training data
    model.fit(scale_train, y_train)

    # Predict using each individual model
    y_pred = model.predict(scale_test)

    # Calculate accuracy
    accuracy = accuracy_score(y_test, y_pred)

    # Store the accuracy of each classifier
    accuracies[name] = accuracy

# Calculate the accuracy for the Voting Classifier
voting_accuracy = accuracy_score(y_test, vpred)
accuracies['Voting Classifier'] = voting_accuracy

# Print out the accuracy of all models
print(accuracies)
```

{'Logistic Regression': 0.843486266687126, 'KNeighbor Classifier': 0.8635875402792696, 'Voting Classifier': 0.8683443302132883}

```
In [114]: import pandas as pd

# Save to a DataFrame for easier handling
accuracies_df = pd.DataFrame(list(accuracies.items()), columns=['Model', 'Accuracy'])

# Save to a CSV file
accuracies_df.to_csv('model_accuracies.csv', index=False)
```

```
# Optionally, print the accuracies DataFrame  
print(accuracies_df)
```

```
Model Accuracy  
0 Logistic Regression 0.843486  
1 KNeighbor Classifier 0.863588  
2 Voting Classifier 0.868344
```

Model Evaluation

```
In [115]: # Evaluate Voting Classifier accuracy  
vaccuracy = accuracy_score(vpred, y_test)  
  
print("Voting Classifier Accuracy:", vaccuracy)
```

```
Voting Classifier Accuracy: 0.8683443302132883
```

```
In [116]: print(len(y_test))  
print(len(vpred))  
# Print the classification report for the Voting Classifier  
evaluation_report = classification_report(vpred, y_test, output_dict=True)
```

```
6517  
6517
```

```
In [117]: # Convert to DataFrame for saving  
report_before_df = pd.DataFrame(evaluation_report).transpose()  
report_before_df.to_csv("report.csv")
```

Check For Class Imbalance

```
In [119]: # Check for class imbalance  
def check_class_imbalance(data, column):  
    class_count = data[column].value_counts()  
    class_frequencies = class_count / class_count.sum()  
    print("Class Frequencies")  
    print(class_frequencies)  
    # Save the class frequencies to a DataFrame  
    class_frequencies_df = pd.DataFrame(class_frequencies).reset_index()  
    class_frequencies_df.columns = [column, 'Frequency'] # Rename columns  
  
    # Save to CSV  
    class_frequencies_df.to_csv("class_imbalance.csv", index=False)  
  
check_class_imbalance(df, 'loan_status')
```

```
Class Frequencies  
0    0.781836  
1    0.218164  
Name: loan_status, dtype: float64
```

Handle Class Imbalance

```
In [120]: # Handle class imbalance using SMOTE  
from imblearn.over_sampling import SMOTE  
  
smote = SMOTE(random_state=42)  
X_resampled, y_resampled = smote.fit_resample(X,y)  
  
# Split the resampled data into training and testing sets  
X_train, X_test, y_train, y_test = train_test_split(X_resampled, y_resampled, test_size=0.2, random_state=5)  
  
# Standardize the resampled feature variables  
X_train_transform = scale.fit_transform(X_train)  
X_test_transform = scale.transform(X_test)  
  
# Fit the Voting Classifier model on the resampled data  
vc.fit(X_train_transform, y_train)  
vpred_2 = vc.predict(X_test_transform)
```

```
In [121]: print(len(y_test))
print(len(vpred))
# Print the classification report for the Voting Classifier
evaluation_report_after = classification_report(vpred_2, y_test, output_dict=True)
# Convert to DataFrame for saving
report_after_df = pd.DataFrame(evaluation_report_after).transpose()
report_after_df.to_csv("report_after.csv")
```

10190
6517

```
In [122]: # Generate confusion matrix and save it
conf_matrix = confusion_matrix(y_test, vpred_2)
conf_matrix_df = pd.DataFrame(conf_matrix, columns=['Predicted 0', 'Predicted 1'], index=['Actual 0', 'Actual 1'])

# Save confusion matrix to CSV
conf_matrix_df.to_csv('confusion_matrix.csv')
```

Cross Validation

In []:

In [34]:

```
# Perform cross-validation
cv_scores = cross_val_score(vc, X_train_transform, y_train, cv=5) # You can adjust the number of folds (cv) as per requirement

# Print the cross-validation scores
print("Cross-validation scores:", cv_scores)
print("Mean CV score:", cv_scores.mean())
```

Cross-validation scores: [0.83918057 0.83216783 0.83388541 0.83155441 0.8340081]
Mean CV score: 0.834159264370772

```
In [124]: # Save the cross-validation scores to a DataFrame
cv_scores_df = pd.DataFrame(cv_scores, columns=['CV_Score'])

# Save to CSV
cv_scores_df.to_csv('cross_validation_scores.csv', index=False)
```

```
In [37]: # Save accuracy result into a dictionary or DataFrame
accuracy_results = pd.DataFrame({
    'Model': ['Voting Classifier'],
    'Accuracy': [vaccuracy]
})

# Save it to CSV
accuracy_results.to_csv('model_accuracy.csv', index=False)
```

```
In [38]: df = pd.read_csv("model_accuracy.csv")
```

```
In [39]: df.head()
```

Out [39]:

	Model	Accuracy
0	Voting Classifier	0.868344

```
In [ ]: import os
import joblib

# Saving the model
model_dir = r'C:\Users\VALERIE KELECHUKWU\Documents\Work_Projects\CProject\RiskApp\RiskAppModel'
joblib.dump(vc, os.path.join(model_dir, 'RiskApp2.6.pkl'))
```

```
In [ ]: print(f"Model saved at: {os.path.join(model_dir, 'CRisk.pkl')}")
```

```
In [ ]: import xgboost  
print(xgboost.__version__)
```

In []:

In []: